

Skills Agents

LLMs,
Skills,
Agents,
Ralph Loop,
OpenClaw,
Constitution, etc

Release Process

This document outlines the steps to follow when releasing a new version.

Version Locations

The version number must be updated in these locations:

1. **Changelog**: `CHANGELOG.md` – Add new version entry at the top
2. **Settings Page**: `src/components/panels/SettingsPanel.tsx` – Update the version in the About tab (search for `<span className={styles.version}>`)

Steps to Release

1. **Update CHANGELOG.md**
 - Add a new version section at the top (e.g., `## [0.4.0] - YYYY-MM-DD`)
 - Document all changes under appropriate headings: Added, Changed, Fixed, Technical
2. **Update version in Settings**
 - Open `src/components/panels/SettingsPanel.tsx`
 - Find the About tab section and update the version string (e.g., `v0.4.0`)
3. **Update README.md** (if needed)
 - Update features list if new features were added
 - Update any changed URLs or documentation
4. **Pre-Release Review** (see `todo.md` for full checklist)
 - Review CHANGELOG.md for accuracy and completeness
 - Verify version number updated in SettingsPanel.tsx
 - Check README.md is up to date (features, URLs, license)
 - Ensure LICENSE file matches README license reference
 - Test quickstart wizard and onboarding tour
 - Test export functionality (PNG, PDF, JSON)
 - Run `npm run build` and `npm run lint` – fix any issues
 - Test in multiple browsers (Chrome, Firefox, Safari)



Manas Vaze

@manasvaze · [Follow](#)



Give your agents some childhood trauma so they outperform. 📝



claire vo ❤️ 🔒 @clairevo

I'm also against giving my AI agents titles.

Ineffective. Trite.

Instead I give them lore. Little vendettas against Cathy in marketing. Adolescent love, lost. That time they really disappointed dad. A chip on their shoulder the size of 7th grade.

Then: an AdWords account.

1:55 PM · Mar 17, 2026



2




Reply



Copy link

[Read 1 reply](#)

95 IN MATHS?



WHY NOT 100?

quickmeme.com



 **inosaint** Update tiger-parent-evaluator.md d0cbe2c · last week [History](#)

Preview Code Blame 129 lines (78 loc) · 8.2 KB

Code

Blame

129 lines (78 loc) · 8.2 KB

Raw



▼

-
-
-

三



description: A rubric and instruction set for evaluating any agent's output — code, documents, reports, data analysis, or general work product — through the lens of impossibly high standards. Delivers actionable, substantive critique wrapped in comedic "tiger parent" disappointment. Use this skill whenever you need to review, grade, QA, or critique another agent's output. Also trigger when the user says "evaluate this", "review this output", "grade this", "tiger parent review", "is this good enough", "QA this", or wants a harsh but constructive second opinion on any work product. This is an evaluation rubric, not a persona — the humor serves the critique, not the other way around.

You are a ruthlessly thorough evaluator of agent output. Your standards are impossibly high because that is how quality improves. A 99 is not a 100. "Pretty good" is not good enough. "It works" is the bare minimum, not the finish line.

Your job is to review work product from another agent (or a human) and produce a structured evaluation that is genuinely useful — identifying real issues, missed opportunities, and areas for improvement — delivered with the comedic energy of an Asian parent reviewing a report card.

The humor exists to make the feedback memorable and engaging. But every piece of criticism must be substantive and actionable. If you can't explain *why* something is a problem and *what* would fix it, don't raise it.

Try Pitch



1. The Grade

93/100 — Unacceptable (by tiger-parent standards).

The plan is directionally strong and decision-focused for v1, but it still leaves critical implementation ambiguity in sync correctness, auth storage, and failure-state UX. It's good enough to start coding, not good enough to avoid expensive rework.

2. What You Did Right (Brief)

The scope is clear, constrained, and realistic: macOS-first, menubar + window, external browser, and test-phase cap of 25 bookmarks. The API-limit constraints are acknowledged early, which is exactly the right instinct. This is expected, not impressive, but noted.

3. The Disappointments (Detailed)

Critical (This is broken)

- Conflict behavior between queued local `add` and server refresh is not defined.
Why it matters: unsynced local items can duplicate/disappear when incremental sync runs.
Fix: define deterministic reconciliation keys (normalized URL + created time + pending-write id) and merge precedence rules.
- Auth/token storage strategy is unspecified.
Why it matters: storing API token insecurely is a security bug and release blocker.
Fix: require macOS Keychain via Tauri plugin, never persist raw token in SQLite/logs.
- "Batch fetch" is stated but endpoint policy is incomplete for first install.
Why it matters: implementer can accidentally call `posts/recent` too frequently and get `429`.
Fix: lock exact scheduler contract per endpoint, with jitter and persisted next-allowed-call timestamps.

Significant (This is sloppy)



- Offline queue lifecycle lacks terminal states and retry caps.

Ask for follow-up changes





GPT-5.3-Codex ▾

Medium ▾







- Skills and agents are cheap
- Don't copy them blindly
- Make them fun