

# Plan: Generate BLR2B Speaker Application PDF

## Goal

Produce a single-page PDF ( `blr2b-application.pdf` ) containing the speaker application email body below, formatted cleanly so it reads as a polished document — not a screenshot of an email.

## Instructions for Claude

1. **Read the pdf skill first.** Before writing any code, view `/mnt/skills/public/pdf/SKILL.md` . Use `reportlab` with `SimpleDocTemplate` and `Paragraph` objects (Platypus flow), not raw `canvas.drawString` . Platypus handles wrapping and spacing; canvas does not.
2. **Use only built-in fonts** (Helvetica family). No Unicode subscripts/superscripts. Em-dashes (—) are fine.
3. **Page setup**
  - Letter size
  - Margins: 1 inch (72pt) on all sides
  - Single page — the content fits if you respect the spacing below
4. **Typography**
  - Subject line: Helvetica-Bold, 12pt, with the label "Subject:" in regular weight before it
  - Body: Helvetica, 11pt, leading 15pt
  - Bullets: same body style, indented, with a hanging em-dash as the bullet marker (not the default round bullet)
  - 8pt spacer between paragraphs; 14pt spacer before and after the bullet block
5. **Content** — use exactly this text. Do not paraphrase, do not "improve" it.

**Subject:** BLR2B speaker application — Skills vs Sub-agents, demystified

**Body:**

Hi Vikram,

Applying to speak at BLR2B tomorrow. My focus: Skills and Sub-agents — specifically, when to reach for which, and why most people get this wrong.

The AI space is buried in jargon — LLMs, Skills, Agents, Ralph Loop, OpenClaw, Constitution, Identity, Soul. The vocabulary makes it feel larger than it is. An “agent” is just an instance of the model with a defined role and context. A skill is a function. A constitution is a decision-making prompt. A soul is a personality profile. All text files. Once you see this, the field gets a lot smaller and a lot more usable.

What I’d cover in 10 minutes:

— Why LLMs default to generic output, and what a skill actually changes — The trap of using a skill you don’t understand (e.g., a motion design skill is useless if you don’t know motion design vocabulary) — My rule of thumb: skills for one-time tasks (metricize-recipe, event-curation), sub-agents for re-usable flows (releases, implementations) — A concrete example from my coffee grind size calibrator: two sub-agents for adding a new grinder — Opus for research planning, Sonnet for implementation. If these were skills, both would run on whatever model the main chat is using. Sub-agents let you be token-conservative and match model strength to task difficulty.

I’ve given an earlier talk on skills and agents, so this isn’t my first time framing it for a room. The angle for BLR2B would be sharper: not a tour of features, but the mental model that lets you stop guessing.

Happy to be in the lineup tomorrow if it’s useful.

Thanks, [Your name]

## 6. Output

- Save the final file to `/mnt/user-data/outputs/blr2b-application.pdf`
- Then call `present_files` with that path

## Quality checks before finishing

- Open the generated PDF and verify it is exactly one page
- Verify no text is clipped at the bottom margin
- Verify the bullet block reads as a list (visible indent + hanging dash), not as a paragraph
- Verify “[Your name]” is present at the bottom so the user knows to replace it before

sending

## **Do not**

- Do not add a header, logo, footer, page number, or any “designed” embellishment. It’s an email rendered as a document, not a brochure.
- Do not split into multiple pages.
- Do not include links, URLs, or QR codes.
- Do not change the wording.